

HAO WANG

Herndon, VA | 215-266-4475 | hwang_pa@msn.com | [LinkedIn](#)

Summary

Senior Software Engineer with extensive experience building system orchestration platforms and integrating AI/LLM capabilities into production systems. Top 2% Cursor AI user pioneering AI-assisted development at Walmart. Deep expertise in building platforms that coordinate multiple systems to solve customer problems - exactly what AI agents need. Strong track record of taking ownership in ambiguous environments and leading teams through unknowns.

Strong in: Python, Go, AI/LLM integration, system orchestration, AWS, Kubernetes, React

Skills

Languages: Python, Go, Java, JavaScript/TypeScript, SQL

AI/ML: LLM APIs (OpenAI, Claude), AI-assisted development, Cursor AI (top 2%), prompt engineering, conversational interfaces

Backend: AWS, Kubernetes, microservices, REST APIs, event-driven architectures

Data Systems: Kafka, SQL Server, MongoDB, Elasticsearch, Memcached

Frontend: React.js, modern web development

Monitoring: Splunk, Grafana, distributed tracing, real-time alerting

Experience

Staff Technical Expert

Walmart Global Tech | Herndon, VA | March 2020 – Present

Building platforms that orchestrate multiple systems to solve problems at scale. Pioneering AI/LLM adoption across the team while leading technical strategy and mentoring engineers.

AI/LLM Integration & Innovation:

AI-Assisted Development Pioneer — Top 2% Cursor AI user globally

- First on team to adopt AI-assisted development, achieving 40% productivity increase
- Integrated LLM APIs (OpenAI, Claude) into internal tooling for automated analysis

- Built AI-powered report generation reducing manual work by 60%
- Evangelized AI adoption across engineering organization
- Mentored engineers on prompt engineering and AI workflows
- This experience prepared me to think about how AI agents orchestrate systems

System Orchestration Platforms (Core AI Agent Pattern):

Workflow Manager & Team Portal — Platform orchestrating 10+ internal systems

- React.js frontend, Python/Go backend serving 50+ engineers
- Orchestrates: data access, job scheduling, secrets management, deployment pipelines
- **This is the AI agent pattern:** coordinate multiple systems to solve user problems
- Self-service interface abstracting complexity - users state intent, platform figures out execution
- Reduced time-to-production by 60% through intelligent orchestration
- Built conversational-style UI guiding users through complex workflows

Event Automation Platform — Multi-system orchestration engine

- Go, GraphQL, React.js, REST APIs, Slack integration
- Orchestrates deployments across 4 regions, multiple services, various approval systems
- **AI agent parallel:** takes high-level intent ("deploy to production"), orchestrates all necessary steps
- Auto-pilot mode for scheduled operations
- Real-time status updates and notifications (conversational feedback loop)
- Handles 100+ orchestrated workflows per week

Backend Systems at Scale:

Real-Time Data Pipeline — Event-driven microservices

- Python/Go services processing millions of events/second
- AWS infrastructure, Kafka messaging, REST APIs
- 99.99% uptime SLA - built for reliability
- Comprehensive monitoring and alerting

Account Risk API — Low-latency microservice

- Node.js/Python API serving real-time decisions
- Sub-50ms p99 latency, millions of requests daily
- Event-driven architecture with Kafka backbone
- Experience building customer-facing APIs

Multi-Signal Scoring System — Coordinating multiple data sources

- Orchestrates external APIs (MaxMind, Spur) and internal services
- Real-time feature aggregation from diverse sources
- **Similar to AI agents:** gather signals from multiple systems, make decision, take action
- Handles Black Friday traffic (10x normal) without issues

Infrastructure & DevOps:

Kubernetes Migration (WCNP) — Leading containerization effort

- Migrating services to Kubernetes, building deployment automation
- Docker, K8s, Helm charts, monitoring
- Building self-service platform for container deployments

Leadership & Team Development:

- Mentor 8 engineers on technical skills and AI-assisted workflows
- Run design reviews, set coding standards, drive technical direction
- Lead cross-functional initiatives across security, product, and platform teams
- On-call rotation: incident response, post-mortems, continuous improvement
- Built culture of ownership and quality through leading by example

Principal Software Engineer / Engineering Lead

Law School Admission Council (LSAC) | Newtown, PA | November 1996 – March 2020

Built production systems ranging from ML applications to distributed platforms.

ML/AI Production Systems:

Recommendation Engine — Personalized matching system

- Python, Pandas, collaborative filtering
- Microservices architecture for serving recommendations
- Early experience with ML systems in production

OCR Pipeline — Document intelligence system

- Python, Tesseract, computer vision
- Production ML pipeline for document processing

Platform Systems:

Digital LSAT Platform — Distributed coordination system (US Patent 10,078,968)

- Coordinated thousands of simultaneous users across platforms
- Real-time orchestration at scale

SSO Platform — Identity system with Azure AD B2C

- Microservices, REST APIs, user behavior analytics

Data Sync Services — Real-time integration pipelines

- Orchestrated data flow across Dynamics 365, Oracle, SQL Server, Salesforce
-

IT Consultant

Fortune 500 Companies | Early Career

Enterprise software development across multiple industries.

Education

Master of Science, Computer Science

Saint Joseph's University — Philadelphia, PA

Patent

System and Method for Electronic Test Delivery

US Patent 10,078,968 | [View on Google Patents](#)

Why Affirm AI Agents?

I've spent the last 4 years building platforms that orchestrate multiple systems to solve problems - exactly what AI agents do. I'm a top 2% AI adopter who's integrated LLMs into production tools. I thrive in ambiguous, newly formed teams where ownership and execution matter most. This role combines my system orchestration expertise with my AI adoption journey - it's exactly where I want to go next.

What I bring:

- Proven AI/LLM integration experience in production
- Deep system orchestration skills (coordinating multiple systems to solve problems)
- Backend expertise: Python, AWS, Kubernetes, microservices at scale
- Frontend skills: React.js for building user interfaces
- Leadership in ambiguous environments - I've launched new initiatives from scratch
- Mentorship and team development experience

Compensation target: \$210K-\$240K (top of range for CA/WA/NY/NJ/CT)